

The NCBI C++ Toolkit

31: Library Configuration

Last Update: June 20, 2013.

Overview

The overview for this chapter consists of the following topics:

- [Introduction](#)
- [Chapter Outline](#)

Introduction

This chapter describes the run-time configuration parameters of the NCBI C++ Toolkit libraries. Such parameters change the default behavior of applications built using the Toolkit.

Configuration parameters can be set by environment variables, entered into a configuration file, defined by code, or any combination of those methods. **Note:** If a parameter is specified in both a configuration file and the environment, the environment takes precedence. The methods supported by each library and application are described below.

Chapter Outline

The following is an outline of the topics presented in this chapter:

- [Defining and Using Parameters](#)
 - [CParam](#)
 - [Registry](#)
 - [Environment](#)
- [Non-Specific Parameters](#)
 - [Logging](#)
 - [Diagnostic Trace](#)
 - [Run-Time](#)
 - [Abnormal Program Termination](#)
 - [NCBI](#)
- [Library-Specific Parameters](#)
 - [Connection](#)
 - [CGI and FCGI](#)
 - [Serial](#)
 - [Objects, Object Manager, Object Tools](#)
 - [cSRA](#)
 - ◆ [sra_read library](#)
 - ◆ [ncbi_xloader_csra library](#)
 - [DBAPI](#)
 - [Eutils](#)
- [Application-Specific Parameters](#)

- [NetCache and NetSchedule](#)
- [Seqfetch.cgi](#)

Defining and Using Parameters

The following sections discuss the methods that libraries can use to define configuration parameters, and the corresponding methods that client applications can use to specify values for those parameters.

- [CParam](#)
- [Registry](#)
- [Environment](#)

CParam

Note: The preferred way for libraries to define their configuration parameters is with the macros in the CParam class (e.g. `NCBI_PARAM_DECL`). More details on the CParam class and its macros are presented in an earlier chapter. Libraries that use CParam can get configuration parameters using either the registry or the environment. Also, the CParam value can be stored and accessed on different levels: globally (application wide) and/or per-thread (TLS-like) and/or locally (cached within a CParam instance). Note that the name of an environment variable linked to a CParam can be customized or follow the default naming convention, so you have to look up the actual name used in the tables below before setting a configuration parameter using the environment.

Registry

If the CParam class cannot be used, the registry (configuration file) may be used to load, access, modify and store the values read from a configuration file. For libraries that use the registry, client applications can set the library configuration parameters using either the registry or the environment. In these cases the environment variable must follow the default naming convention.

These environment variables can be used to specify where to look for the registry.

The registry is case-insensitive for section and entry names. More details on the registry are presented in an earlier chapter.

Environment

For configuration parameters defined by either CParam or the registry, there is an equivalent environment variable having the form `NCBI_CONFIG__<section>__<name>` (note the double-underscores preceding `<section>` and `<name>`). The equivalent form is all uppercase.

Note: Environment variables may not contain dots (a.k.a. period or full stop) on many platforms. However, dots are allowed in registry section and entry names. The equivalent environment variable for parameters containing a dot in the section or entry name is formed by replacing the period with `_DOT_`. For example, the equivalent environment variable for [FastCGI]

WatchFile.Name is `NCBI_CONFIG__FASTCGI__WATCHFILE_DOT_NAME`.

Note: Environment variables are case-sensitive on many platforms. Therefore, when setting a configuration parameter via the environment, be sure to use the case shown in the tables below.

Some configuration parameters can only be set with an environment variable - for example, DIAG_SILENT_ABORT. In such cases, there is no corresponding registry entry.

Non-Specific Parameters

The following sections discuss configuration parameters that are not library-specific.

- [Logging](#)
- [Diagnostic Trace](#)
- [Run-Time](#)
- [Abnormal Program Termination](#)
- [NCBI](#)

Logging

The application log consists of diagnostic messages. Some of them are available only in debug builds. Others - namely, those produced by the ERR_POST or LOG_POST macros - can be redirected into a file. Normally, the name and location of the application log is specified using the logfile command-line argument.

These parameters tune the usage and behavior of the application log file.

Diagnostic Trace

These parameters tune the visibility and contents of diagnostic messages produced by _TRACE or ERR_POST macros.

See Table 3.

Run-Time

Run-time configuration parameters allow specifying memory size limit, CPU time limit, and memory allocation behavior. **Note:** not all operating systems support these parameters.

Abnormal Program Termination

These parameters specify how to handle abnormal situations when executing a program.

NCBI

These parameters tune generic NCBI C++ Toolkit-wide behavior.

Library-Specific Parameters

The following sections discuss library-specific configuration parameters.

- [Connection](#)
- [NetCache and NetSchedule](#)
- [CGI and FCGI](#)
- [Serial](#)
- [Objects, Object Manager, Object Tools](#)
- [DBAPI](#)
- [Eutils](#)

Connection

These parameters affect various aspects of internet connections established by the connection library. See the Networking and IPC chapter for a description of the corresponding network information structure.

CGI and FCGI

These parameters tune the behavior of CGI and FCGI applications and built with the NCBI C++ Toolkit libraries. See Table 10 for CGI Load balancing configuration parameters.

Serial

These parameters tune the behavior of the Serial library.

Objects, Object Manager, Object Tools

These parameters tune the behavior of the Objects-related libraries, including the Object Manager and loader and reader libraries.

cSRA

sraread library

Note: This section applies only inside NCBI.

The following parameters tune the behavior of the sraread library:

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|--|---|--------------|---------|
| If true, will add CIGAR info to Seq-align's returned by cSRA iterators. | [csra] cigar_in_align_ext CSRA_CIGAR_IN_ALIGN_EXT | Boolean | true |
| If true, will clip the read ranges returned by cSRA short read iterators according to quality. | [csra] clip_by_quality CSRA_CLIP_BY_QUALITY | Boolean | true |
| If true, will add mate info to Seq-align's returned by cSRA iterators. | [csra] explicit_mate_info CSRA_EXPLICIT_MATE_INFO | Boolean | false |
| If true, cSRA short read iterators will also include technical reads. | [csra] include_technical_reads CSRA_INCLUDE_TECHNICAL_READS | Boolean | true |

ncbi_xloader_csra library

Note: This section applies only inside NCBI.

The following parameters tune the behavior of the ncbi_xloader_csra library:

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|--|---|--------------|---------|
| If ≥ 9 , log alignment chunks. If ≥ 5 , log major function calls. If ≥ 2 , log refseq stats. If ≥ 1 , log summary data. | [csra_loader] debug CSRA_LOADER_DEBUG | int | 0 |
| The max number of SRR files to keep open. | [csra_loader] gc_size CSRA_LOADER_GC_SIZE | size_t | 10 |
| If > 0 , defines the max number of separate spot groups. | [csra_loader] max_separate_spot_groups CSRA_LOADER_MAX_SEPARATE_SPOT_GROUPS | int | 0 |
| If > 0 , defines the minimum quality threshold for loading alignment and pileup chunks. | [csra_loader] pileup_graphs CSRA_LOADER_PILEUP_GRAPHs | int | 0 |
| If true, fetch quality graphs along with short reads. | [csra_loader] quality_graphs CSRA_LOADER_QUALITY_GRAPHs | Boolean | false |

DBAPI

These parameters tune the behavior of the DBAPI library.

Eutils

These parameters tune the behavior of the Eutils library.

Application-Specific Parameters

The following sections discuss configuration parameters that are specific to selected applications.

- [NetCache and NetSchedule](#)
- [Seqfetch.cgi](#)

NetCache and NetSchedule

Note: This applies only inside NCBI.

Table 16 describes configuration parameters that are common to both NetCache and NetSchedule client APIs. These parameters are found in the netservice_api registry section.

Table 17 describes configuration parameters for NetCache client applications. These parameters are found in the netcache_api registry section. **Note:** The netcache_api registry section was formerly called netcache_client.

Table 18 describes configuration parameters for NetSchedule client applications. These parameters are found in the netschedule_api registry section.

See the Distributed Computing chapter for more information on NetCache and NetSchedule.

Seqfetch.cgi

Note: This applies only inside NCBI.

These parameters tune the behavior of the seqfetch.cgi application.

Table 1. Registry configuration parameters

| Purpose | Environment variable | Valid values |
|--|----------------------------|--------------|
| If this variable is defined, the value is an extra-high-priority configuration file whose entries override those from other configuration files. | NCBI_CONFIG_OVERRIDES | a valid path |
| If this variable is defined, use it exclusively as the registry search path. | NCBI_CONFIG_PATH | a valid path |
| If this variable is not defined, append the current directory and home directory to the registry search path (after NCBI_CONFIG_PATH). | NCBI_DONT_USE_LOCAL_CONFIG | anything |
| If this variable is defined, append the value to the registry search path (after the home directory). | NCBI | a valid path |
| For Windows: If this variable is defined, append the value to the registry search path (after NCBI). For non-Windows, this variable is not checked and /etc is appended to the registry search path (after NCBI). | SYSTEMROOT | a valid path |
| If this variable is not defined, attempt to load a low-priority system-wide registry (ncbi.ini on Windows; .ncbirc on non-Windows). Note: the system-wide registry will not be loaded if it contains the DONT_USE_NCBIRC entry in the NCBI section. | NCBI_DONT_USE_NCBIRC | anything |

Table 2. Log file configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|--|--|---|------------------|
| Used by logging framework if the real client IP can not be obtained. | [LOG] Client_Ip NCBI_LOG_CLIENT_IP | a valid IPv4 or IPv6 address | "" |
| Reset the log file to the specified file. | [LOG] File NCBI_CONFIG__LOG__FILE ^c | a valid file name | "" |
| Specify when to use the File, NoCreate, Truncate, and TryRootLogFirst registry parameters shown in this table. Note: those parameters will only be used if the log file has not been set already or if IgnoreEnvArg is set to true. | [LOG] IgnoreEnvArg NCBI_CONFIG__LOG__IGNOREENVARG ^c | Boolean ^a | false |
| The listed environment variables will be logged as an 'extra' after each 'request-start' message. The extra message starts with a "LogEnvironment=true" pair. | [LOG] LogEnvironment DIAG_LOG_ENVIRONMENT [sic] | space separated list of environment variable names | "" |
| The listed registry entries will be logged as an 'extra' after each 'request-start' message. The extra message starts with a "LogRegistry=true" pair. | [LOG] LogRegistry DIAG_LOG_REGISTRY [sic] | space separated list of registry section:name values | "" |
| Do not create the log file if it does not exist already. | [Log] NoCreate NCBI_CONFIG__LOG__NOCREATE ^c | Boolean ^b | false |
| Specifies what to do if an invalid page hit ID is encountered. Valid PHIDs match the regex /[A-Za-z0-9:@_]+(\.[0-9]+)*/. | [Log] On_Bad_Hit_Id LOG_ON_BAD_HIT_ID | "Allow", "AllowAndReport", "Ignore", "IgnoreAndReport", "Throw" | "AllowAndReport" |
| Specifies what to do if an invalid session ID is encountered. Valid session IDs match the format specified by LOG_SESSION_ID_FORMAT. | [Log] On_Bad_Session_Id LOG_ON_BAD_SESSION_ID | "Allow", "AllowAndReport", "Ignore", "IgnoreAndReport", "Throw" | "AllowAndReport" |
| Turn performance logging on or off (globally). | [Log] PerfLogging LOG_PerfLogging ^c | Boolean ^b | false |
| Defines the default session ID, which is used for any request which has no explicit session ID set. | [Log] Session_Id NCBI_LOG_SESSION_ID | any valid session ID string | "" |

| | | | |
|---|---|-----------------------------|------------|
| Specifies which format rule to check session IDs against: for "Ncbi" use <code>^[0-9]{16}_[0-9]{4,}SID\$</code> for "Standard" use <code>^[A-Za-z0-9_\.:@-]+\$</code> for "Other" use <code>^.*\$</code> (i.e. anything is valid). | [Log] Session_Id_Format LOG_SESSION_ID_FORMAT | "Ncbi", "Standard", "Other" | "Standard" |
| If set, a "log_site" entry with the given value will be added to request-start entries in log files. | [LOG] Site NCBI_LOG_SITE | a URL-encoded site name | (none) |
| If this parameter is defined, use the CSysLog facility setting when posting. | [LOG] SysLogFacility NCBI_CONFIG__LOG__SYSLOGFACILITY ^c | any non-empty string | (none) |
| Truncate the log file – i.e. discard the contents when opening an existing file. | [Log] Truncate LOG_TRUNCATE | Boolean ^b | false |
| Specify whether to try creating the log file under /log before trying other locations (e.g. a location specified by the registry or by NCBI_CONFIG__LOG__FILE). | [LOG] TryRootLogFirst NCBI_CONFIG__LOG__TRYROOTLOGFIRST ^c | Boolean ^a | false |
| If true, default to logging warnings when unsafe static array types are copied. | [NCBI] STATIC_ARRAY_COPY_WARNING NCBI_STATIC_ARRAY_COPY_WARNING | Boolean ^b | false |
| If true, log warnings for unsafe static array types. | [NCBI] STATIC_ARRAY_UNSAFE_TYPE_WARNING NCBI_STATIC_ARRAY_UNSAFE_TYPE_WARNING | Boolean ^b | true |

^a case-insensitive: true, t, yes, y, false, f, no, n

^b case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^c environment variable name formed from registry section and entry name

Table 3. Diagnostic trace configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|--|---|--|---------|
| Specify the severity level threshold for posting diagnostic messages – i.e. less severe messages will not be posted. Note: If the parameter is set then the function <code>ncbi::SetDiagPostLevel()</code> is disabled - except for setting the level to <code>eDiag_Trace</code> . | [DEBUG] DIAG_POST_LEVEL DIAG_POST_LEVEL | CI ^b : Info, Warning, Error, Critical, Fatal, Trace | (none) |
| Diagnostic trace will be enabled if this parameter is given any value. | [DEBUG] DIAG_TRACE DIAG_TRACE or NCBI_CONFIG__DEBUG__DIAG_TRACE ^c | any non-empty string | (none) |
| Specify a file that stores a mapping of error codes to their descriptions. | [DEBUG] MessageFile NCBI_CONFIG__DEBUG__MessageFile ^c | a valid file name | (none) |
| Specify the maximum number of messages that can be posted to the AppLog within the AppLog period. | [Diag] AppLog_Rate_Limit DIAG_APPLOG_RATE_LIMIT | unsigned integer | 50000 |
| Specify the AppLog period in seconds. | [Diag] AppLog_Rate_Period DIAG_APPLOG_RATE_PERIOD | unsigned integer | 10 |
| Specify whether context properties should be automatically printed when set or changed. | [Diag] AutoWrite_Context DIAG_AUTOWRITE_CONTEXT | Boolean ^a | false |
| Specify the maximum number of diagnostic messages to collect. Messages beyond the limit will result in erasing the oldest message. | [Diag] Collect_Limit DIAG_COLLECT_LIMIT | size_t | 1000 |
| Specify the maximum number of messages that can be posted to the ErrLog within the ErrLog period. | [Diag] ErrLog_Rate_Limit DIAG_ERRLOG_RATE_LIMIT | unsigned integer | 5000 |
| Specify the ErrLog period in seconds. | [Diag] ErrLog_Rate_Period DIAG_ERRLOG_RATE_PERIOD | unsigned integer | 1 |
| Limit the log file size, and rotate the log when it reaches the limit. | [Diag] Log_Size_Limit DIAG_LOG_SIZE_LIMIT | non-negative long integer | 0 |
| Use the old output format if the flag is set. | [Diag] Old_Post_Format DIAG_OLD_POST_FORMAT | Boolean ^a | true |

| | | | |
|---|--|--|----------------------------------|
| Specify a diagnostics post filter string (see an earlier chapter for more detail on filtering). | [DIAG] POST_FILTER NCBI_CONFIG_DIAG_POST_FILTER ^c | see the syntax rules | (none) |
| Print the system TID rather than CThread::GetSelf(). | [Diag] Print_System_TID DIAG_PRINT_SYSTEM_TID | Boolean ^a | false |
| Defines the maximum number of entries to be listed in a stack trace. All stack trace entries above the specified level are not printed. | [DIAG] Stack_Trace_Max_Depth DEBUG_STACK_TRACE_MAX_DEPTH | a positive integer | 200 |
| Specify the minimum severity that will activate Tee_To_Stderr. See the Tee Output to STDERR section. | [Diag] Tee_Min_Severity DIAG_TEE_MIN_SEVERITY | CI ^b : Info, Warning, Error, Critical, Fatal, Trace | Warning (debug); Error (release) |
| Duplicate messages to stderr. See the Tee Output to STDERR section. | [Diag] Tee_To_Stderr DIAG_TEE_TO_STDERR | Boolean ^a | false |
| Specify a diagnostics trace filter string (see an earlier chapter for more detail on filtering). | [DIAG] TRACE_FILTER NCBI_CONFIG_DIAG_TRACE_FILTER ^c | see the syntax rules | (none) |
| Specify the maximum number of messages that can be posted to the TraceLog within the TraceLog period. | [Diag] TraceLog_Rate_Limit DIAG_TRACELOG_RATE_LIMIT | unsigned integer | 5000 |
| Specify the TraceLog period in seconds. | [Diag] TraceLog_Rate_Period DIAG_TRACELOG_RATE_PERIOD | unsigned integer | 1 |
| If true and AppLog severity is not locked, print the current GMT time in diagnostic messages; otherwise print local time. | [Diag] UTC_Timestamp DIAG_UTC_TIMESTAMP | Boolean ^a | false |

^a case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^b CI = case-insensitive

^c environment variable name formed from registry section and entry name

Table 4. Run-time configuration parameters

| Purpose | [Registry section] Registry name | Valid values | Default |
|---|--|--|---------------|
| | Environment variable | | |
| Set a CPU time limit for the application in seconds. | [NCBI] CpuTimeLimit NCBI_CONFIG__NCBI__CPUTIMELIMIT ^b | non-negative integer | 0 (unlimited) |
| Set a memory size limit for the application in MB or as a percent of total system memory. | [NCBI] MemorySizeLimit NCBI_CONFIG__NCBI__MEMORYSIZELIMIT ^b | non-negative integer or percent, for example "70" (for 70 MB) or "70%" (for 70% of all memory) | 0 (unlimited) |
| Specify the method for filling allocated memory. | [NCBI] MEMORY_FILL NCBI_MEMORY_FILL | CI ^a : none, zero, pattern | pattern |

^a CI = case-insensitive

^b environment variable name formed from registry section and entry name

Table 5. Abnormal program termination configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|---|--|----------------------|---------|
| If this parameter is defined, abort the program if a CException is thrown. | [DEBUG] ABORT_ON_THROW NCBI_CONFIG__DEBUG__ABORT_ON_THROW ^c | any non-empty string | (none) |
| Specify whether the NCBI application framework should catch exceptions that are not otherwise caught. | [Debug] Catch_Unhandled_Exceptions DEBUG_CATCH_UNHANDLED_EXCEPTIONS | Boolean ^a | true |
| Specify whether ncbi::Abort() will call _ASSERT (false). Note: this only applies to MSVC. | [Diag] Assert_On_Abort DIAG_ASSERT_ON_ABORT | Boolean ^a | false |
| If this parameter is true, abort the program if a CObjectException is thrown. | [NCBI] ABORT_ON_COBJECT_THROW NCBI_ABORT_ON_COBJECT_THROW | Boolean ^a | false |
| If this parameter is true, abort the program on an attempt to access or release a NULL pointer stored in a CRef object. | [NCBI] ABORT_ON_NULL NCBI_ABORT_ON_NULL | Boolean ^a | false |
| Specify what to do when ncbi::Abort() is called. When the variable is set to a "yes" value, Abort() will call exit(255). When the variable is set to a "no" value, Abort() will call abort(). When the variable is not set, Abort() will call exit(255) for release builds and abort() for debug builds - unless compiled with MSVC and the DIAG_ASSERT_ON_ABORT parameter is true, in which case Abort() will call _ASSERT(false). | [N/A] N/A DIAG_SILENT_ABORT | Boolean ^b | (none) |

^a case-insensitive: true, t, yes, y, 1, false, f, no, n, 0^b case-insensitive: y, 1, n, 0^c environment variable name formed from registry section and entry name

Table 6. NCBI C++ Toolkit-wide configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|---|--|--|-----------------------------------|
| Specify whether throwing an exception of at least Critical severity will cause an immediate abort(). | [EXCEPTION] Abort_If_Critical EXCEPTION_ABORT_IF_CRITICAL | Boolean ^a | false |
| Specify the minimum severity that will result in the stack trace being added to exceptions. | [EXCEPTION] Stack_Trace_Level EXCEPTION_STACK_TRACE_LEVEL | CI ^b : Trace, Info, Warning, Error, Critical, Fatal | Critical |
| A single path to check for common data files via g_FindDataFile(). Takes a lower precedence than paths in NCBI_DATA_PATH. | [NCBI] Data NCBI_CONFIG__NCBI__DATA ^c | a valid path | "" |
| A list of paths (delimited in the style of the OS) to check for common data files via g_FindDataFile(). | [NCBI] DataPath NCBI_DATA_PATH | a delimited list of valid paths | "" |
| Specify how read-only files are treated on Windows during a remove request. | [NCBI] DeleteReadOnlyFiles NCBI_CONFIG__DELETEREADONLYFILES | Boolean ^a | false |
| Specify whether the API classes should have logging turned on. | [NCBI] FileAPILogging NCBI_CONFIG__FILEAPILOGGING | Boolean ^a | DEFAULT_LOGGING_VALUE |
| Declare how umask settings on Unix affect creating files/ directories in the File API. | [NCBI] FileAPIHonorUmask NCBI_CONFIG__FileAPIHonorUmask | Boolean ^a | false |
| Specify whether to load plugins from DLLs. | [NCBI] Load_Plugins_From_DLLs NCBI_LOAD_PLUGINS_FROM_DLLS | Boolean ^a | LOAD_PLUGINS_FROM_DLLS_BY_DEFAULT |
| Specify the directory to use for temporary files. | [NCBI] TmpDir NCBI_CONFIG__NCBI__TMPDIR ^c | a valid path | "" |
| Specify the file name of a Unicode-to-ASCII translation table. | [NCBI] UnicodeToAscii NCBI_CONFIG__NCBI__UNICODETOASCII ^c | a valid path | "" |

^a case-insensitive: true, t, yes, y, 1, false, f, no, n, 0^b CI = case-insensitive

^c environment variable name formed from registry section and entry name

Table 7. Connection library configuration parameters

| Purpose | [Registry section] Registry name Environment variable (See <i>Note 2</i>) | Valid values | Default |
|--|--|---|--------------------|
| Service-specific parameters follow this form. (See <i>Note 1</i>) | [<service>] CONN_<param_name> <service>_CONN_<param_name> | | |
| Global parameters follow this form. (See <i>Note 1</i>) | [CONN] <param_name> CONN_<param_name> | | |
| Specify arguments for the given service. (See <i>Note 1</i>) | [<service>] CONN_ARGS <service>_CONN_ARGS | (service-dependent) | "" |
| Specify how much debug information will be output. (See <i>Note 1</i>) | [<service>] CONN_DEBUG_PRINTOUT <service>_CONN_DEBUG_PRINTOUT | CI ^a : to get some: 1, on, yes, true, some to get all: data, all to get none: anything else | "" |
| If this parameter is true, the network dispatcher will be disabled. (See <i>Note 1</i>) | [<service>] CONN_DISPD_DISABLE <service>_CONN_DISPD_DISABLE | Boolean ^c | true |
| If this parameter is true, the Firewall mode will be enabled. (See <i>Note 1</i>) | [<service>] CONN_FIREWALL <service>_CONN_FIREWALL | Boolean ^c | not set |
| Set the dispatcher host name. (See <i>Note 1</i>) | [<service>] CONN_HOST <service>_CONN_HOST | a valid host name | www.ncbi.nlm.nih.g |
| Set the HTTP proxy server. (See <i>Note 1</i>) | [<service>] CONN_HTTP_PROXY_HOST <service>_CONN_HTTP_PROXY_HOST | a valid proxy host | "" |
| Set the HTTP proxy server port number. This will be set to zero if <service>_CONN_HTTP_PROXY_HOST is not set. (See <i>Note 1</i>) | [<service>] CONN_HTTP_PROXY_PORT <service>_CONN_HTTP_PROXY_PORT | unsigned short | 0 |
| Set a custom user header. This is rarely used, and then typically for debugging purposes. (See <i>Note 1</i>) | [<service>] CONN_HTTP_USER_HEADER <service>_CONN_HTTP_USER_HEADER | a valid HTTP header | "" |
| Prohibit the use of a local load balancer. <i>Note</i> : This parameter is discouraged for performance reasons - please use <service>_CONN_LBSMD_DISABLE instead. (See <i>Note 1</i>) | [<service>] CONN_LB_DISABLE <service>_CONN_LB_DISABLE | Boolean ^c | false |

| | | | |
|---|---|------------------------------------|--------------------|
| Prohibit the use of a local load balancer. This should be used instead of <code><service> CONN_LB_DISABLE</code> . (See <i>Note 7</i>) | [<service>] CONN_LBSMD_DISABLE <service>_CONN_LBSMD_DISABLE | Boolean ^c | false |
| Enable the use of locally configured services. See <code><service> CONN_LOCAL_SERVER_<n></code> . (See <i>Note 7</i>) | [<service>] CONN_LOCAL_ENABLE <service>_CONN_LOCAL_ENABLE | Boolean ^c | false |
| Create a service entry for service, where n is a number from 0 to 100 (not necessarily sequential). The value must be a valid server descriptor, as it would be configured for the load balancing daemon (LBSMD). This is a quick way of configuring locally used services (usually, for the sole purposes of debugging / development) without the need to edit the actual LBSMD tables (which become visible for the whole NCBI). See <code><service> CONN_LOCAL_ENABLE</code> . Note: This parameter has no corresponding global parameter. (See <i>Note 1</i>) | [<service>] CONN_LOCAL_SERVER_<n> <service>_CONN_LOCAL_SERVER_<n> | any non-empty string | not set |
| Maximum number of attempts to establish connection. Zero means use the default. (See <i>Note 1</i>) | [<service>] CONN_MAX_TRY <service>_CONN_MAX_TRY | unsigned short | 3 |
| Specify a password for the connection (only used with <code><service>_CONN_USER</code>). (See <i>Note 1</i>) | [<service>] CONN_PASS <service>_CONN_PASS | the user's password | "" |
| Set the path to the service. (See <i>Note 1</i>) | [<service>] CONN_PATH <service>_CONN_PATH | a valid service path | /Service/dispd.cgi |
| Set the dispatcher port number. (See <i>Note 1</i>) | [<service>] CONN_PORT <service>_CONN_PORT | unsigned short | 0 |
| Set a non-transparent CERN-like firewall proxy server. (See <i>Note 1</i>) | [<service>] CONN_PROXY_HOST <service>_CONN_PROXY_HOST | a valid proxy host | "" |
| Set the HTTP request method. (See <i>Note 1</i>) | [<service>] CONN_REQ_METHOD <service>_CONN_REQ_METHOD | CI ^a : any, get, post | ANY |
| Redirect connections to <service> to the specified alternative service. See Service Redirection. (See <i>Note 1</i>) | [<service>] CONN_SERVICE_NAME <service>_CONN_SERVICE_NAME | a replacement for the service name | (none) |
| Set to true if the client is stateless. (See <i>Note 1</i>) | [<service>] CONN_STATELESS <service>_CONN_STATELESS | Boolean ^c | false |

| | | | |
|---|--|--|--|
| Zero means no waiting but polling (may not work well with all connections); "infinite" means no timeout (i.e. to wait for I/O indefinitely); other values are the maximum number of seconds to wait before failing. (See <i>Note 1</i> .) | [<service>] CONN_TIMEOUT <service>_CONN_TIMEOUT | floating point >= 0.0 (1 microsecond precision) ^f or "infinite" | 30.0 |
| Specify a username for the connection (see <service>_CONN_PASS). Only necessary for connections requiring authentication. (See <i>Note 1</i> .) | [<service>] CONN_USER <service>_CONN_USER | a username with access rights for the connection | (none) |
| Set the level of logging detail that GNUTLS should produce about secure transactions. Log levels greater than 7 also dump scrambled data from GNUTLS. | [CONN] GNUTLS_LOGLEVEL CONN_GNUTLS_LOGLEVEL | 0 to 10 | 0 |
| A true value enables HTTP connections to dump headers of error server responses only (successful responses do not get logged). | [CONN] HTTP_ERROR_HEADER_ONLY CONN_HTTP_ERROR_HEADER_ONLY | Boolean ^c | false |
| A true value enables HTTP connections to follow https to http transitions (http to https transitions are secure and therefore don't need to be enabled). | [CONN] HTTP_INSECURE_REDIRECT CONN_HTTP_INSECURE_REDIRECT | Boolean ^c | false |
| Set a default referer (applies to all HTTP connections). | [CONN] HTTP_REFERER CONN_HTTP_REFERER | a valid referer | (none) |
| A list of identifiers to be treated as local services defined in the registry / environment. This parameter is optional and is used only for reverse address-to-name lookups. | [CONN] LOCAL_SERVICES CONN_LOCAL_SERVICES | whitespace-delimited ^d list of identifiers | (none) |
| Set the mail gateway host. | [CONN] MX_HOST CONN_MX_HOST | a valid host name | localhost on UNIX platforms except Cygwin; mailgw on other platforms |
| Set the mail gateway port. | [CONN] MX_PORT CONN_MX_PORT | 1 to 65535 | 25 (SMTP) |
| Set the mail gateway communication timeout in seconds. | [CONN] MX_TIMEOUT CONN_MX_TIMEOUT | floating point >= 0.0 (zero means default) | 120 |
| Enable CServer to catch exceptions. | [server] Catch_Unhandled_Exceptions CSERVER_CATCH_UNHANDLED_EXCEPTIONS | Boolean ^b | true |
| Deprecated. | [server] allow_implicit_job_return NCBI_CONFIG_SERVER_ALLOW_IMPLICIT_JOB_RETURN ^e | Boolean ^b | false |
| Maximum time worker nodes are allowed to live without a single NetSchedule server. | [server] max_wait_for_servers NCBI_CONFIG_SERVER_MAX_WAIT_FOR_SERVERS ^e | unsigned int | 24 * 60 * 60 seconds |

| | | | |
|---|--|----------------------|------|
| Causes the worker node to shut down if any jobs fail. | [server] stop_on_job_errors NCBI_CONFIG__SERVER__STOP_ON_JOB_ERRORS ^e | Boolean ^b | true |
| Enable CThreadInPool_ForServer to catch exceptions. | [ThreadPool] Catch_Unhandled_Exceptions NCBI_CONFIG__THREADPOOL__CATCH_UNHANDLED_EXCEPTIONS ^e | Boolean ^b | true |

^a CI = case-insensitive

^b case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^c case-insensitive: true values are { 1, on, yes, true }; false is anything else

^d whitespace can be any number of spaces and/or tabs

^e environment variable name formed from registry section and entry name

^f although very precise values may be specified, practical host limitations may result in less precise effective values

Note 1: All service-specific parameters shown in Table 7 (except one) have corresponding global parameters - i.e. parameters that apply to all services. For these global parameters, the registry section name is CONN; the registry entry name doesn't have the CONN_ prefix; and the environment variable doesn't have the <service>_ prefix. For example, the service-specific parameter specified by the CONN_ARGS entry in a given [<service>] section of the registry (or by the <service>_CONN_ARGS environment variable) corresponds to the global parameter specified by the ARGS entry in the [CONN] section of the registry (or by the CONN_ARGS environment variable). When both a service-specific parameter and its corresponding global parameter are set, the service-specific parameter takes precedence.

Note 2: Environment variable names for service-specific parameters are formed by capitalizing the service name.

Table 8. CGI-related configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|---|--|---|----------|
| Set to the user agent string you would like to be used by the web server. | [N/A] N/A HTTP_USER_AGENT | A valid user agent string. | (none) |
| Add to the user agent list of bot names. This parameter affect only CCGiUserAgent::IsBot(). | [CGI] Bots NCBI_CONFIG_CGI_BOTS ^f | Delimited list ^b of bot names, e.g. "Googlebot Scooter WebCrawler Slurp". | (none) |
| According to RFC-2109, cookies should not be encoded. Instead, they should be just quoted. However, for backward compatibility with code that decodes incoming cookies, both quoted cookies and encoded cookies can be parsed. This setting controls which method of encoding/decoding is used. | [CGI] Cookie_Encoding CGI_COOKIE_ENCODING | "Url", "Quote" | "Url" |
| Severity level for cookie-related error messages. | [CGI] Cookie_Error_Severity CGI_Cookie_Error_Severity | CI ^e : Info, Warning, Error, Critical, Fatal, Trace | Error |
| Defines which characters cannot be used in cookie names. | [CGI] Cookie_Name_Banned_Symbols CGI_Cookie_Name_Banned_Symbols | A string of banned characters. | " ,;=" |
| Set to true to make the application count the amount of data read/sent. The numbers are then printed in request stop log messages. | [CGI] Count_Transfered CGI_COUNT_TRANSFERED | Boolean ^c | true |
| Set the name of an environment variable, which in turn specifies a prefix that will be added to all diagnostic messages issued during HTTP request processing. | [CGI] DiagPrefixEnv NCBI_CONFIG_CGI_DIAGPREFIXENV ^f | a valid environment variable name | (none) |
| Set to true to disable the creation of a tracking cookie during session initialization. | [CGI] DisableTrackingCookie NCBI_CONFIG_CGI_DISABLETRACKINGCOOKIE ^f | Boolean ^c | false |
| Set to true to enable logging. | [CGI] Log NCBI_CONFIG_CGI_LOG ^f | CI ^e : On => enabled; True => enabled; OnError => enabled for errors; OnDebug => enabled (debug builds only) | disabled |
| An ampersand-delimited string of GET and/or POST arguments to exclude from the log (helps limit the size of the log file) | [CGI] LOG_EXCLUDE_ARGS CGI_LOG_EXCLUDE_ARGS | valid format: arg1&arg2... | (none) |

| | | | |
|---|--|--|-----------|
| Allows specifying limits for multiple GET and/or POST arguments in one parameter string. | [CGI] LOG_LIMIT_ARGS CGI_LOG_LIMIT_ARGS | valid format: arg1:size1&arg2:size2...&*:size special argument: * means all unspecified arguments; special limits: -2 means exclude; -1 means no limit | *:1000000 |
| Enable logging of CGI request parameters. Only the specified parameters will be logged. | [CGI] LogArgs NCBI_CONFIG_CGI_LOGARGS ^f | Delimited list ^b of environment variables (optionally aliased on output for shortening logs, e.g. envvar=1). | (none) |
| Set to true to merge log lines. | [CGI] Merge_Log_Lines CGI_MERGE_LOG_LINES | Boolean ^c | true |
| Specify additional mobile device names. This parameter affect only CCgiUserAgent::IsMobileDevice(). | [CGI] MobileDevices NCBI_CONFIG_CGI_MobileDevices ^f | Delimited list ^b of additional device names. | (none) |
| Add to the user agent list of names that aren't bots. This parameter affect only CCgiUserAgent::IsBot(). | [CGI] NotBots NCBI_CONFIG_CGI_NotBots ^f | Delimited list ^b of names that aren't bots. | (none) |
| Add to the user agent list of names that aren't mobile devices. This parameter affect only CCgiUserAgent::IsMobileDevice(). | [CGI] NotMobileDevices NCBI_CONFIG_CGI_NotMobileDevices ^f | Delimited list ^b of names that aren't mobile devices. | (none) |
| Add to the user agent list of names that aren't phone devices. This parameter affect only CCgiUserAgent::IsPhoneDevice(). | [CGI] NotPhoneDevices NCBI_CONFIG_CGI_NotPhoneDevices ^f | Delimited list ^b of names that aren't phone devices. | (none) |
| Add to the user agent list of names that aren't tablet devices. This parameter affect only CCgiUserAgent::IsTabletDevice(). | [CGI] NotTabletDevices NCBI_CONFIG_CGI_NotTabletDevices ^f | Delimited list ^b of names that aren't tablet devices. | (none) |
| Control error handling of incoming cookies (doesn't affect outgoing cookies set by application). | [CGI] On_Bad_Cookie CGI_ON_BAD_COOKIE | CI ^e : Throw, SkipAndError, Skip, StoreAndError, Store | Store |
| Specify additional phone device names. This parameter affect only CCgiUserAgent::IsPhoneDevice(). | [CGI] PhoneDevices NCBI_CONFIG_CGI_PhoneDevices ^f | Delimited list ^b of additional device names. | (none) |
| Specifies whether to print the referer during LogRequest(). | [CGI] Print_Http_Referer CGI_PRINT_HTTP_REFERER | Boolean ^c | true |
| Specifies whether to print the URL during LogRequest(). | [CGI] Print_Self_Url CGI_PRINT_SELF_URL | Boolean ^c | true |

| | | | |
|---|--|---|--------------|
| Specifies whether to print the user agent during LogRequest(). | [CGI] Print_User_Agent CGI_PRINT_USER_AGENT | Boolean ^c | true |
| Set the size of CGI request buffer that is printed when the request cannot be parsed. | [CGI] RequestErrBufSize NCBI_CONFIG_CGI_REQUESTERRBUFSIZE ^f | buffer size in bytes | 256 |
| Specify the registry section name for the result cache. | [CGI] ResultCacheSectionName NCBI_CONFIG_CGI_RESULTCACHESECTIONNAME ^f | valid section name | result_cache |
| Enable statistics logging. | [CGI] StatLog NCBI_CONFIG_CGI_STATLOG ^f | Boolean ^d | false |
| Specify additional tablet device names. This parameter affect only CCGiUserAgent::IsTabletDevice(). | [CGI] TabletDevices NCBI_CONFIG_CGI_TabletDevices ^f | Delimited list ^b of additional device names. | (none) |
| Controls whether the output stream will throw for bad states. | [CGI] ThrowOnBadOutput NCBI_CONFIG_CGI_THROWONBADOUTPUT ^f | Boolean ^c | true |
| Log start time, end time, and elapsed time. | [CGI] TimeStamp NCBI_CONFIG_CGI_TIMESTAMP ^f | Boolean ^d | false |
| Disable statistics logging if the CGI request took less than the specified number of seconds. | [CGI] TimeStatCutOff NCBI_CONFIG_CGI_TIMESTATCUTOFF ^f | non-negative integer (zero enables logging) | 0 |
| Specify the domain for the tracking cookie. | [CGI] TrackingCookieDomain NCBI_CONFIG_CGI_TRACKINGCOOKIEDOMAIN ^f | valid domain | .nih.gov |
| Specify the tracking cookie name. | [CGI] TrackingCookieName NCBI_CONFIG_CGI_TRACKINGCOOKIE ^f | valid cookie name | ncbi_sid |
| Specify the path for the tracking cookie. | [CGI] TrackingCookiePath NCBI_CONFIG_CGI_TRACKINGCOOKIEPATH ^f | valid path | / |
| Defines the name of the NCBI tracking cookie (session ID cookie). | [CGI] TrackingTagName CGI_TrackingTagName | Any valid cookie name. | "NCBI-SID" |

^a List may be delimited by semicolon, space, tab, or comma.

^b List may be delimited by semicolon, space, tab, vertical bar, or tilde.

^c case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^d case-insensitive: true, t, yes, y, false, f, no, n

^e CI = case-insensitive

^f environment variable name formed from registry section and entry name

Table 9. FCGI-related configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|--|--|---|---------|
| A true value enables logging of current iteration, max iterations, and process ID during the FastCGI run. | [FastCGI] Debug NCBI_CONFIG_FASTCGI_DEBUG ^b | Boolean ^a | false |
| A true value enables termination of a FastCGI application by the presence of the request entry "exitfastcgi". | [FastCGI] HonorExitRequest NCBI_CONFIG_FASTCGI_HONOREXITREQUEST ^b | Boolean ^a | false |
| Specify the number of requests that the FCGI application will process before exiting. | [FastCGI] Iterations NCBI_CONFIG_FASTCGI_ITERATIONS ^b | positive integer | 10 |
| Make the FastCGI application run as a stand-alone server on a local port. The value is a UNIX domain socket or a MS Windows named pipe, or a colon followed by a port number | [FastCGI] StandaloneServer FCGI_STANDALONE_SERVER | valid local port or named socket | (none) |
| Make the FastCGI application stop if an error is encountered. | [FastCGI] StopIfFailed NCBI_CONFIG_FASTCGI_STOPIFFAILED ^b | Boolean ^a | false |
| Make the FastCGI application exit if the named file changes. | [FastCGI] WatchFile.Name NCBI_CONFIG_FASTCGI_WATCHFILE_DOT_NAME ^b | valid file name | (none) |
| The number of bytes to read from the watch file to see if it has changed. | [FastCGI] WatchFile.Limit NCBI_CONFIG_FASTCGI_WATCHFILE_DOT_LIMIT ^b | positive integer (non-positives trigger default) | 1024 |
| The period in seconds between checking the watch file for changes. | [FastCGI] WatchFile.Timeout NCBI_CONFIG_FASTCGI_WATCHFILE_DOT_TIMEOUT ^b | positive integer (non-positives trigger default, which is to disable the watch file checking) | 0 |

^a case-insensitive: true, t, yes, y, false, f, no, n

^b environment variable name formed from registry section and entry name

Table 10. CGI Load balancing configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|---|--|----------------------|-------------------|
| Specify the internet domain. | [CGI-LB] Domain NCBI_CONFIG__CGI-LB__DOMAIN ^b | a valid domain | .ncbi.nlm.nih.gov |
| Specify the host IP address. | [CGI-LB] Host NCBI_CONFIG__CGI-LB__HOST ^b | a valid host IP | (none) |
| Specify the cookie expiration period in seconds. | [CGI-LB] LifeSpan NCBI_CONFIG__CGI-LB__LIFESPAN ^b | integer | 0 |
| Specify the name of the load balancing cookie in the HTTP response. | [CGI-LB] Name NCBI_CONFIG__CGI-LB__NAME ^b | a valid cookie name | (none) |
| Specify the cookie path. | [CGI-LB] Path NCBI_CONFIG__CGI-LB__PATH ^b | a valid path | (none) |
| Specify the cookie security mode. | [CGI-LB] Secure NCBI_CONFIG__CGI-LB__SECURE ^b | Boolean ^a | false |

^a case-insensitive: true, t, yes, y, false, f, no, n

^b environment variable name formed from registry section and entry name

Table 11. Serial library configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|---|--|--|--------------------|
| Skip unknown data members in the input stream, or throw an exception. | [N/A] N/A SERIAL_SKIP_UNKNOWN_MEMBERS | CI ^a : yes, no, never, always | no (throw) |
| If true, causes CObjectOutputStream::WriteDouble () to use fast conversion. | [SERIAL] FastWriteDouble NCBI_CONFIG_SERIAL_FastWriteDouble ^b | Boolean ^c | true |
| While reading binary ASN.1 data allow VisibleString tag where UTF-8 string tag is expected by specification. | [SERIAL] READ_ANY_UTF8STRING_TAG SERIAL_READ_ANY_UTF8STRING_TAG | Boolean ^c | true |
| While reading binary ASN.1 data allow UTF-8 string tag where VisibleString tag is expected by specification. | [SERIAL] READ_ANY_VISIBLESTRING_TAG SERIAL_READ_ANY_VISIBLESTRING_TAG | 0 (disallow, throws an exception); 1 (allow, but warn once); 2 (allow without warning) | 1 |
| Specify how to handle unknown variants when reading Object streams. | [SERIAL] SKIP_UNKNOWN_MEMBERS NCBI_CONFIG_SERIAL_SKIP_UNKNOWN_MEMBERS ^b | CI ^a : no (throw an exception), never (even if set to skip later), yes (skip), always (even if set to not skip later) | no |
| Specify how to handle unknown variants when reading Object streams. | [SERIAL] SKIP_UNKNOWN_VARIANTS NCBI_CONFIG_SERIAL_SKIP_UNKNOWN_VARIANTS ^b | CI ^a : no (throw an exception), never (even if set to skip later), yes (skip), always (even if set to not skip later) | no |
| Throw an exception on an attempt to access an uninitialized data member. | [SERIAL] VERIFY_DATA_GET SERIAL_VERIFY_DATA_GET | CI ^a : yes, no, never, always, defvalue, defvalue_always | yes |
| Throw an exception if a mandatory data member is missing in the input stream. | [SERIAL] VERIFY_DATA_READ SERIAL_VERIFY_DATA_READ | CI ^a : yes, no, never, always, defvalue, defvalue_always | yes |
| Throw an exception on an attempt to write an uninitialized data member. | [SERIAL] VERIFY_DATA_WRITE SERIAL_VERIFY_DATA_WRITE | CI ^a : yes, no, never, always, defvalue, defvalue_always | yes |
| While writing binary ASN.1 data issue UTF8 string tag as determined by specification, otherwise issue plain string tag. | [SERIAL] WRITE_UTF8STRING_TAG SERIAL_WRITE_UTF8STRING_TAG | Boolean ^c | false |
| Specifies what to do if an invalid character is read. | [SERIAL] WRONG_CHARS_READ NCBI_CONFIG_SERIAL_WRONG_CHARS_READ ^b | "ALLOW", "REPLACE", "REPLACE_AND_WARN", "THROW", "ABORT" | "REPLACE_AND_WARN" |

| | | | |
|--|--|--|--------------------|
| Specifies what to do if an invalid character is written. | [SERIAL] WRONG_CHARS_WRITE NCBI_CONFIG_SERIAL_WRONG_CHARS_WRITE ^b | "ALLOW", "REPLACE", "REPLACE_AND_WARN", "THROW", "ABORT" | "REPLACE_AND_WARN" |
|--|--|--|--------------------|

^a CI = case-insensitive

^b environment variable name formed from registry section and entry name

^c case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

Table 13. Objects-related configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|--|--|---|------------------|
| The Object Manager will attach WGS master descriptors to Bioseq data by default. Setting this parameter to false will disable this behavior. | [GENBANK] ADD_WGS_MASTER GENBANK_ADD_WGS_MASTER | Boolean ^a | true |
| A non-zero value turns on debugging messages about GenBank loader's interaction with cache. | [GENBANK] CACHE_DEBUG GENBANK_CACHE_DEBUG | >=0, currently only zero and non-zero are distinguished | 0 |
| Specify whether an attempt should be made to recompress the cache. | [GENBANK] CACHE_RECOMPRESS GENBANK_CACHE_RECOMPRESS | Boolean ^a | true |
| A non-zero value turns on debugging messages about opening/closing connections to ID1/ID2 services. | [GENBANK] CONN_DEBUG GENBANK_CONN_DEBUG | >=0, currently only zero and non-zero are distinguished | 0 |
| Disable attaching WGS master descriptors when retrieving ASN.1 blobs using the CPubseqReader and CPubseq2Reader classes. | [GENBANK/PUBSEQOS] or [GENBANK/PUBSEQOS2] EXCLUDE_WGS_MASTER NCBI_CONFIG_GENBANK_PUBSEQOS_EXCLUDE_WGS_MASTER or NCBI_CONFIG_GENBANK_PUBSEQOS2_EXCLUDE_WGS_MASTER | Boolean ^b | false |
| Set the severity level for ID1 debug tracing. | [GENBANK] ID1_DEBUG GENBANK_ID1_DEBUG | int: 0 = none, 1 = error, 2 = open, 4 = conn, 5 = asn, 8 = asn data | 0 |
| Specify the ID1 reader service name. Note: The services can be redirected using generic Service Redirection technique. | In priority order: [GENBANK] ID1_SERVICE_NAME, [NCBI] SERVICE_NAME_ID1 In priority order: GENBANK_ID1_SERVICE_NAME, GENBANK_SERVICE_NAME_ID1 | a valid reader service name | ID1 (see API) |
| Specify the ID2 reader service name. Note: The services can be redirected using generic Service Redirection technique. | In priority order: [GENBANK] ID2_CGI_NAME, [GENBANK] ID2_SERVICE_NAME, [NCBI] SERVICE_NAME_ID2 In priority order: GENBANK_ID2_CGI_NAME, GENBANK_ID2_SERVICE_NAME, GENBANK_SERVICE_NAME_ID2 | a valid reader service name | ID2 (see API) |

| | | | |
|---|--|---|---|
| Set the severity level for ID2 debug tracing. | [GENBANK] ID2_DEBUG GENBANK_ID2_DEBUG | int: 0 = none, 1 = error, 2 = open, 4 = conn, 5 = asn, 8 = blob, 9 = blob data | debug: none release: error (see API) |
| Number of chunks allowed in a single request. | [GENBANK] ID2_MAX_CHUNKS_REQUEST_SIZE GENBANK_ID2_MAX_CHUNKS_REQUEST_SIZE | int: 0 = unlimited request size; 1 = do not use packets or get-chunks requests | 100 |
| Maximum number of requests packed in a single ID2 packet. | [GENBANK] ID2_MAX_IDS_REQUEST_SIZE GENBANK_ID2_MAX_IDS_REQUEST_SIZE | >=0 | 100 |
| The maximum number of connections the reader can establish to the data source. This is run-time limited to 1 for single threaded clients and for all clients using the cache or gi reader, and to 5 for multi-threaded clients using the id1, id2, pubseqos, and pubseqos2 readers. | [GENBANK] MAX_NUMBER_OF_CONNECTIONS | int | 3 for id1 and id2; 2 for pubseqos and pubseqos2 |
| See MAX_NUMBER_OF_CONNECTIONS | [GENBANK] NO_CONN | | |
| See OPEN_TIMEOUT_INCREMENT | [GENBANK] OPEN_INCREMENT | | |
| See OPEN_TIMEOUT_MAX | [GENBANK] OPEN_MAX | | |
| See OPEN_TIMEOUT_MULTIPLIER | [GENBANK] OPEN_MULTIPLIER | | |
| The OPEN_TIMEOUT* parameters describe the timeout for opening a GenBank connection. The timeout allows the server a reasonable time to respond while providing a means to quickly abandon unresponsive servers. | [GENBANK] OPEN_TIMEOUT NCBI_CONFIG_GENBANK_OPEN_TIMEOUT ^c | any floating point value >= 0.0 | 5 seconds |
| OPEN_TIMEOUT_MULTIPLIER and OPEN_TIMEOUT_INCREMENT specify the way the open timeout is increased if no response is received (next_open_timeout = prev_open_timeout * multiplier + increment). | [GENBANK] OPEN_TIMEOUT_INCREMENT NCBI_CONFIG_GENBANK_OPEN_TIMEOUT_INCREMENT ^c | any floating point value >= 0.0 | 0 seconds |
| The limit of increasing the open timeout using OPEN_TIMEOUT_MULTIPLIER and OPEN_TIMEOUT_INCREMENT. | [GENBANK] OPEN_TIMEOUT_MAX NCBI_CONFIG_GENBANK_OPEN_TIMEOUT_MAX ^c | floating point >= 0.0 | 30 seconds |
| See OPEN_TIMEOUT_INCREMENT | [GENBANK] OPEN_TIMEOUT_MULTIPLIER NCBI_CONFIG_GENBANK_OPEN_TIMEOUT_MULTIPLIER ^c | floating point >= 0.0 | 1.5 |

| | | | |
|---|---|--|---|
| Turns on different levels of debug messages in PubSeqOS reader. A value ≥ 2 means debug opening connections while ≥ 5 means debug results of Seq-id resolution requests. Note: only applies to debug builds. | [GENBANK] PUBSEQOS_DEBUG GENBANK_PUBSEQOS_DEBUG | int | 0 |
| Whether to open first connection immediately or not. | [GENBANK] preopen NCBI_CONFIG__GENBANK__PREOPEN ^c | Boolean ^b | true |
| Specify the level of reader statistics to collect. | [GENBANK] READER_STATS GENBANK_READER_STATS | int: 0 = none, 1 = verbose | 0 |
| Prioritized list of drivers to try for the reader. | Sources searched for list: [GENBANK] ReaderName, [GENBANK] LOADER_METHOD, default Sources searched for list: GENBANK_LOADER_METHOD, default | list items are semicolon- delimited; each item is a colon- delimited list of drivers. valid drivers: id1, id2, cache, pubseqos | "ID2:PUBSEQOS:ID1", or "ID2:ID1" (see API) |
| Specify whether the reader manager should automatically register ID1, ID2, and cache. | [GENBANK] REGISTER_READERS GENBANK_REGISTER_READERS | Boolean ^a | true |
| Specify whether the blob stream processor should try to use string packing. | [N/A] N/A NCBI_SERIAL_PACK_STRINGS | Boolean ^d | true |
| On some platforms, equal strings can share their character data, reducing the required memory. Set this parameter to true to have the GenBank loader try to use this feature if it is available. | [GENBANK] SNP_PACK_STRINGS GENBANK_SNP_PACK_STRINGS | Boolean ^a | true |
| In ID1/PubSeqOS readers present SNP data as ID2-split entries to reduce memory usage. | [GENBANK] SNP_SPLIT GENBANK_SNP_SPLIT | Boolean ^a | true |
| Storing all the SNPs as plain ASN.1 objects would require a huge amount of memory. The SNP table is a compact way of storing SNPs to reduce memory consumption. Set this parameter to true to have the object manager try to use the SNP table. | [GENBANK] SNP_TABLE GENBANK_SNP_TABLE | Boolean ^a | true |
| Set to a positive integer to enable dumping (to stderr in text ASN.1 form) all the SNPs that don't fit into the SNP table. Note: this is only available in debug mode. | [GENBANK] SNP_TABLE_DUMP GENBANK_SNP_TABLE_DUMP | Boolean ^a | false |
| Set this parameter to true to dump (to stdout) some statistics on the process of storing SNPs into the SNP table. This option may help determine why not all the SNPs could fit in the table. | [GENBANK] SNP_TABLE_STAT GENBANK_SNP_TABLE_STAT | Boolean ^a | false |

| | | | |
|--|---|--|--|
| Specify whether to use a memory pool. | [GENBANK] USE_MEMORY_POOL GENBANK_USE_MEMORY_POOL | Boolean ^a | true |
| The WAIT_TIME* parameters describe the wait time before opening new GenBank connections in case of communication errors. The wait time is necessary to allow network and/or GenBank servers to recover. WAIT_TIME is the initial wait after the first error. See also: GenBank reader configuration. | [GENBANK] WAIT_TIME NCBI_CONFIG__GENBANK__WAIT_TIME ^c | floating point >= 0.0 | 1 second |
| Specifies for how many sequential communication errors the response should be to use wait time, before trying to open a new connection instead. | [GENBANK] WAIT_TIME_ERRORS NCBI_CONFIG__GENBANK__WAIT_TIME_ERRORS ^c | int | 2 errors |
| WAIT_TIME_MULTIPLIER and WAIT_TIME_INCREMENT specify the way wait time is increased if errors continue to happen (next_wait_time = prev_wait_time * multiplier + increment). | [GENBANK] WAIT_TIME_INCREMENT NCBI_CONFIG__GENBANK__WAIT_TIME_INCREMENT ^c | any floating point value >= 0.0 | 1 second |
| The limit of increasing wait time using WAIT_TIME_MULTIPLIER and WAIT_TIME_INCREMENT. | [GENBANK] WAIT_TIME_MAX NCBI_CONFIG__GENBANK__WAIT_TIME_MAX ^c | floating point >= 0.0 | 30 seconds |
| See WAIT_TIME_INCREMENT | [GENBANK] WAIT_TIME_MULTIPLIER NCBI_CONFIG__GENBANK__WAIT_TIME_MULTIPLIER ^c | any floating point value >= 0.0 | 1.5 |
| Prioritized list of drivers to try for the writer. | Sources searched for list: [GENBANK] WriterName, [GENBANK] LOADER_METHOD, default Sources searched for list: GENBANK_LOADER_METHOD, default | list items are semicolon-delimited; each item is a colon-delimited list of drivers. valid drivers: id1, id2, cache, pubseqos | "ID2:PUBSEQOS:ID1", or "ID2:ID1" (see API) |
| If non-zero, reserve Dense-seg vectors using predefined pre-read hook. | [OBJECTS] DENSE_SEG_RESERVE OBJECTS_DENSE_SEG_RESERVE | Boolean ^a | true |
| If non-zero, reserve Seq-graph vectors using predefined pre-read hook. | [OBJECTS] SEQ_GRAPH_RESERVE OBJECTS_SEQ_GRAPH_RESERVE | Boolean ^a | true |
| If non-zero, reserve Seq-table vectors using predefined pre-read hook. | [OBJECTS] SEQ_TABLE_RESERVE OBJECTS_SEQ_TABLE_RESERVE | Boolean ^a | true |
| Specify whether Seq-id general trees are packed. | [OBJECTS] PACK_GENERAL OBJECTS_PACK_GENERAL | Boolean ^a | true |

| | | | |
|--|--|----------------------|-------|
| Specify whether Seq-id text-seq trees are packed. | [OBJECTS] PACK_TEXTID OBJECTS_PACK_TEXTID | Boolean ^a | true |
| Specify whether empty Seq-descr's will be allowed (or throw if not). | [OBJECTS] SEQ_DESCR_ALLOW_EMPTY OBJECTS_SEQ_DESCR_ALLOW_EMPTY | Boolean ^a | false |
| Sets the maximum number of master TSE blobs that will be cached. | [OBJMGR] BLOB_CACHE OBJMGR_BLOB_CACHE | unsigned int | 10 |
| Specify whether the scope can be auto-released. | [OBJMGR] SCOPE_AUTORELEASE OBJMGR_SCOPE_AUTORELEASE | Boolean ^a | true |
| Specify the size of the scope auto-release. | [OBJMGR] SCOPE_AUTORELEASE_SIZE OBJMGR_SCOPE_AUTORELEASE_SIZE | unsigned int | 10 |
| Specify whether the new FASTA implementation will be used. | [READ_FASTA] USE_NEW_IMPLEMENTATION NCBI_CONFIG_READ_FASTA_USE_NEW_IMPLEMENTATION ^c | Boolean ^a | true |

^a case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^b case-insensitive: true, t, yes, y, false, f, no, n

^c environment variable name formed from registry section and entry name

^d case-insensitive: true values are { yes | 1 }; anything else is false

Table 14. DBAPI configuration parameters

| Purpose | [Registry section] Registry name Environment variable | Valid values | Default |
|---|--|--|---------------|
| If RESET_SYBASE is true, the Sybase client path will be set to the value in the SYBASE variable. | [N/A] N/A RESET_SYBASE | Boolean ^a | (none) |
| If RESET_SYBASE is true, the Sybase client path will be set to the value in the SYBASE variable. | [N/A] N/A SYBASE | a path containing a Sybase client | (none) |
| The version of the TDS protocol to use with the CTLIB driver. | [CTLIB] TDS_VERSION CTLIB_TDS_VERSION | an installed TDS version | 125 (see API) |
| The version of the TDS protocol to use with the FTDS driver. | [FTDS] TDS_VERSION FTDS_TDS_VERSION | 0 (auto-detect), 50 (Sybase or Open Server), 70 (SQL Server) | 0 |
| Whether connecting with Kerberos authentication is supported. If true, and the username and password are empty strings, then DBAPI will attempt to use Kerberos to connect to the database. The user must ensure that the database will allow them to connect via Kerberos and that their Kerberos ticket is not expired. | [dbapi] can_use_kerberos NCBI_CONFIG__DBAPI__CAN_USE_KERBEROS ^c | Boolean ^b | false |
| Whether to encrypt login data. | [dbapi] conn_use_encrypt_data NCBI_CONFIG__DBAPI__CONN_USE_ENCRYPT_DATA ^c | Boolean ^b | false |
| The maximum number of simultaneously open connections to database servers. | [dbapi] max_connection NCBI_CONFIG__DBAPI__MAX_CONNECTION ^c | unsigned int | 100 |

| | | | |
|--|--|--------------|----|
| The maximum number of connection attempts that will be made for any server. | [DB_CONNECTION_FACTORY] MAX_CONN_ATTEMPTS NCBI_CONFIG__DB_CONNECTION_FACTORY__MAX_CONN_ATTEMPTS ^c | unsigned int | 1 |
| The maximum number of validation attempts that will be made for each connection. | [DB_CONNECTION_FACTORY] MAX_VALIDATION_ATTEMPTS NCBI_CONFIG__DB_CONNECTION_FACTORY__MAX_VALIDATION_ATTEMPTS ^c | unsigned int | 1 |
| The maximum number of servers to try to connect to for each service name (this is only meaningful if the number of servers running this service exceeds this value). | [DB_CONNECTION_FACTORY] MAX_SERVER_ALTERNATIVES NCBI_CONFIG__DB_CONNECTION_FACTORY__MAX_SERVER_ALTERNATIVES ^c | unsigned int | 32 |
| The maximum number of connections to be made to one particular server (when several connections to the same service name are requested) before an attempt to connect to another server will be made. A value of 0 means connect to the same server indefinitely. | [DB_CONNECTION_FACTORY] MAX_DISPATCHES NCBI_CONFIG__DB_CONNECTION_FACTORY__MAX_DISPATCHES ^c | unsigned int | 0 |
| The timeout, in seconds, to be used for all connection attempts (0 means to use either the default value or a value set specifically for the driver context). | [DB_CONNECTION_FACTORY] CONNECTION_TIMEOUT NCBI_CONFIG__DB_CONNECTION_FACTORY__CONNECTION_TIMEOUT ^c | unsigned int | 30 |
| The timeout, in seconds, to be used while logging into the server for all connection attempts (0 means to use either the default value or a value set specifically for the driver context). | [DB_CONNECTION_FACTORY] LOGIN_TIMEOUT NCBI_CONFIG__DB_CONNECTION_FACTORY__LOGIN_TIMEOUT ^c | unsigned int | 30 |

| | | | |
|--|--|--|--------------|
| <p>If DBAPI resolved the passed name as a service name and then couldn't connect to any server associated with that service name, then this parameter determines whether DBAPI should also try to resolve the passed name as a server name (a database alias from “interfaces” file or a DNS name). See also: database load balancing.</p> | <p>[DB_CONNECTION_FACTORY] TRY_SERVER_AFTER_SERVICE</p> <p>NCBI_CONFIG__DB_CONNECTION_FACTORY__TRY_SERVER_AFTER_SERVICE ^c</p> | <p>Boolean ^a</p> | <p>false</p> |
| <p>See 'PRAGMA cache_size' in the SQLite documentation.</p> | <p>[LDS2] SQLiteCacheSize</p> <p>LDS2_SQLITE_CACHE_SIZE</p> | <p>any valid cache size for an SQLite database</p> | <p>2000</p> |

^a case-insensitive: true, t, yes, y, false, f, no, n

^b case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^c environment variable name formed from registry section and entry name

Table 15. eutils library configuration parameters

| Purpose | [Registry section] Registry name | Valid values | Default |
|---|-------------------------------------|--------------|---|
| | Environment variable | | |
| Specify the base URL for Eutils requests. | [Eutils] Base_URL | a valid URL | http://eutils.ncbi.nlm.nih.gov/entrez/eutils/ (see API) |
| | EUTILS_BASE_URL | | |

Table 16. Common NetCache and NetSchedule client API configuration parameters (netservice_api)

| Purpose | [Registry section] Registry name Environment variable | Valid va |
|---|--|--|
| Fail the request if the network I/O is inactive (blocked waiting for the communication channel to become readable or writable) for more than the specified timeout in seconds. Applies to all socket operations after the initial connection is established (see NCBI_CONFIG__NETSERVICE_API__CONNECTION_TIMEOUT). Can be overridden by NCBI_CONFIG__NETCACHE_API__COMMUNICATION_TIMEOUT or NCBI_CONFIG__NETSCHEDULE_API__COMMUNICATION_TIMEOUT. | [netservice_api] communication_timeout NCBI_CONFIG__NETSERVICE_API__COMMUNICATION_TIMEOUT ^a | floating point >= (zero me default) |
| The maximum number of times the API will retry a communication command on a socket. Setting connection_max_retries to zero will prevent NetCache API from retrying the connection and command execution | [netservice_api] connection_max_retries NCBI_CONFIG__NETSERVICE_API__CONNECTION_MAX_RETRIES ^a | unsigned |
| The timeout in seconds for establishing a new connection to a server. Can be overridden by NCBI_CONFIG__NETCACHE_API__CONNECTION_TIMEOUT or NCBI_CONFIG__NETSCHEDULE_API__CONNECTION_TIMEOUT. | [netservice_api] connection_timeout N/A | floating point >= millisec precision minimum 0.001 (1 millisec |
| The number of connections to keep in the local connection pool. If zero, the server will grow the connection pool as necessary to accomodate new connections. Otherwise, when all connections in the pool are used, new connections will be created and destroyed. | [netservice_api] max_connection_pool_size NCBI_CONFIG__NETSERVICE_API__MAX_CONNECTION_POOL_SIZE ^a | non- negative |
| The maximum number of attempts to resolve the LBSMD service name. If not resolved within this limit an exception is thrown. | [netservice_api] max_find_laname_retries NCBI_CONFIG__NETSERVICE_API__MAX_FIND_LANAME_RETRIES ^a | positive |
| The delay in seconds between retrying a command; the total time should not exceed NCBI_CONFIG__NETCACHE_API__MAX_CONNECTION_TIME. | [netservice_api] retry_delay NCBI_CONFIG__NETSERVICE_API__RETRY_DELAY ^a | floating point >= |
| Close connections with zero timeout to prevent sockets in TIME_WAIT on the client side. By default, the Linux kernel delays releasing ports for a certain period after close() because there might be a delayed arrival of packets. Setting this parameter to true disables that behavior and therefore allows faster recycling of ports. This is important when the server is handling a large number of connections due to the limited number of ports available. | [netservice_api] use_linger2 NCBI_CONFIG__NETSERVICE_API__USE_LINGER2 ^a | Boolean |

^a environment variable name formed from registry section and entry name

^b case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

Table 17. NetCache client API configuration parameters (netcache_api)

| Purpose | [Registry section] Registry name Environment variable | Valid values |
|--|--|--|
| Enable input caching (provides for slow blob retrieval). | [netcache_api] cache_input N/A | Boolean ^b |
| Only applies when using CNetlCacheClient. Provides a "namespace" for blobs. Thus, blobs are uniquely identified by the { key, version, subkey, cache_name } combination. | [netcache_api] cache_name N/A | up to 36 characters (cas |
| Enable output caching (provides for saving a blob with pauses more than "communication_timeout"). | [netcache_api] cache_output N/A | Boolean ^b |
| The name of your application, as identified to NetCache. | [netcache_api] client N/A | your application's name |
| Synonym for [netcache_api]/client, which is preferred. | [netcache_api] client_name N/A | |
| Can be used to override NCBI_CONFIG__NETSERVICE_API__COMMUNICATION_TIMEOUT. Please see that entry for details. | [netcache_api] communication_timeout N/A | floating point >= 0.0 (z NCBI_CONFIG__NET |
| Can be used to override [netservice_api]/connection_timeout. Please see that entry for details. | [netcache_api] connection_timeout N/A | floating point >= 0.0, m [netservice_api]/conne |
| Depending on the value, enables mirroring: if true, mirroring is unconditionally enabled, if false, it is disabled completely. The special value "if_key_mirrored" is used to enable mirroring for the blobs that already have mirroring extensions in their keys. | [netcache_api] enable_mirroring N/A | Boolean , or "if_key_m |
| The host:port address for the NetCache server that will be used for blob creation if none of the servers configured via LBSM were able to create the blob. This is only for new blob requests. | [netcache_api] fallback_server NCBI_CONFIG__NETCACHE_API__FALLBACK_SERVER ^a | a valid server |
| In conjunction with [netcache_api]/port, a synonym for [netcache_api]/service_name, which is preferred. | [netcache_api] host N/A | |
| Max total time for each NetCache transaction. | [netcache_api] max_connection_time N/A | floating point >= 0.0 (z |
| In conjunction with [netcache_api]/host, a synonym for [netcache_api]/service_name, which is preferred. | [netcache_api] port N/A | |

| | | |
|---|--|---|
| A trigger for LBSM query (query LBSM once per the specified number of NetCache operations). | [netcache_api] rebalance_requests N/A | integer ≥ 0 (zero means no trigger) |
| Another trigger for LBSM query (query LBSM at least once per the specified number of seconds) | [netcache_api] rebalance_time N/A | floating point ≥ 0.0 (zero means no trigger) |
| Synonym for [netcache_api]/host, which is preferred. | [netcache_api] server N/A | |
| Synonym for [netcache_api]/service_name. | [netcache_api] service N/A | |
| The LBSM name that specifies which servers to use. The service name is only used when creating blobs. | [netcache_api] service_name N/A | any registered LBSM service |
| This is one condition that will trigger server throttling and is defined as a string having the form "A / B" where A and B are integers. Throttling will be triggered if there are A failures in the last B operations. | [netcache_api] throttle_by_connection_error_rate N/A | a string having the form "A / B" |
| This is another condition that will trigger server throttling and is defined as follows. Server throttling will be triggered if this number of consecutive connection failures happens. | [netcache_api] throttle_by_consecutive_connection_failures N/A | integer |
| Do not release server throttling until the server appears in LBSMD. | [netcache_api] throttle_hold_until_active_in_lb N/A | Boolean ^c |
| Indicates when server throttling will be released. | [netcache_api] throttle_relaxation_period N/A | integer time period in seconds |
| Where to save blob caches. | [netcache_api] tmp_dir N/A | a valid directory |
| Synonym for [netcache_api]/tmp_dir. | [netcache_api] tmp_path N/A | |
| A true value enables an alternative method for checking if a blob exists. Note: This option is available only for backward compatibility and should not be used. | [netcache_api] use_hasb_fallback NCBI_CONFIG__NETCACHE_API__USE_HASB_FALLBACK ^a | Boolean ^b |
| Defines LBSM affinity name to use for floor assignment, etc. | [netcache_api] use_lbsm_affinity N/A | a valid affinity |

^a environment variable name formed from registry section and entry name

^b case-insensitive: true, t, yes, y, 1, false, f, no, n, 0

^c case-insensitive: true, t, yes, y, false, f, no, n

Table 18. NetSchedule client API configuration parameters (netschedule_api)

| Purpose | [Registry section] Registry name | Valid values |
|--|---|--|
| | Environment variable | |
| Name of the queue (DO NOT use default queue for your application). | [netschedule_api] queue_name N/A | your application's queue name |
| The name of your application, as identified to NetSchedule. | [netschedule_api] client_name N/A | your application's name |
| Can be used to override NCBI_CONFIG__NETSERVICE_API__COMMUNICATION_TIMEOUT. Please see that entry for details. | [netschedule_api] communication_timeout N/A | floating point ≥ 0.0 (zero means use the default from NCBI_CONFIG__NETSERVICE_API__COMMUNICATION_T |
| Can be used to override [netservice_api]/connection_timeout. Please see that entry for details. | [netschedule_api] connection_timeout N/A | floating point ≥ 0.0 (zero means use the default from [netservice_a connection_timeout) |

Table 19. seqfetch.cgi application configuration parameters

| Purpose | [Registry section] Registry name | Valid values | Default |
|---------------------------------------|---|--------------------|----------------------|
| | Environment variable | | |
| Point to the current script. | [SeqFetch] Viewer_fcgi_path SEQFETCH_VIEWER_FCGI_PATH | a valid path | /sviewer/viewer.fcgi |
| Name the current load-balanced proxy. | [SeqFetch] Viewer_fcgi_proxy SEQFETCH_VIEWER_FCGI_PROXY | a valid proxy name | sviewer_lb |